


Client-Server Model

- Usually 2 computers, one provides a service (server), one issues service requests (clients)
- Form of interaction:
 - a *server* starts first and awaits contact
 - a *client* starts second and initiates the connection




- Some examples of client-server model: file transfer service; web service; email service; video on demand service; content sharing service; real-time search service; domain name service; address allocation service..
- Key characters of a **client**
 - client often runs locally on a user's personal computer and is invoked directly by a user
 - client actively initiates contact with the server
 - client can access multiple services as needed (browsers), but usually contacts one remote server at a time
 - client typically does not require powerful computer hardware
- Key characters of a **server**
 - it is dedicated to provide one service that can handle multiple remote clients at the same time
 - it is often invoked automatically when a system boots, and continues to execute through many sessions
 - it waits passively for contact from remote clients
 - it accepts contact from arbitrary clients, but offers a single service
 - it may require powerful hardware and a sophisticated OS
- Server
 - refers to a *program* that waits passively for communication, not the *computer* on which it executes
 - when a computer is dedicated to running one or more service programs, the computer itself is sometimes called a server (ok...)

- Server vs. Client

	Server	Client
Start	Starts first	Starts second
Destination	Does not need to know which client will contact	Must know which server to contact
Connection	Waits passively and arbitrarily long for connection from clients	Initiate a connection whenever a communication is needed
Communication	Communicate with a client by both sending and receiving data	Communicate with a server by both sending and receiving data
Termination	Stays running after serving one client and waits for another	May terminate after interaction with a server

Requests and Responses

- Once contact has been established, two-way communication is possible (i.e., data can flow from a client to a server or from a server to a client)
- In some cases, a client sends a series of requests and the server issues a series of responses (e.g., a database client might allow a user to look up more than one item per session)



Multiple servers

- Allowing a given computer to operate multiple services is helpful because...
 - hardware can be shared
 - single computer has lower system administration overhead than multiple computer systems
 - experience has shown that the demand for a service is often sporadic
 - If demand for a service is low, consolidating services on a single computer can dramatically reduce cost

Multiple clients

- A single computer can run...
 - a single client
 - multiple copies of a client that contact a server (ssh command window * 10 lol)
 - multiple clients that each contact a particular server
- Allowing multiple clients is useful because...
 - servers (services) can be accessed simultaneously
- For example, a user can have three 3 windows open simultaneously running three 3 applications
 - one that retrieves and displays email
 - another that connects to a chat service
 - and a third running a web browser

So we have clients and servers, and we need a formal way to exactly specify the communication between servers and clients; let's consider the **TCP/IP protocol suite**.

TCP/IP protocol suite is a group of protocol and application programming interfaces which are used to communicate over the Internet.

IP: the Internet Protocol - handles **addressing** used to identify specific hosts, and delivery of messages between hosts


TCP: the Transmission Control Protocol - is a **heavyweight** communication protocol which provides a *reliable* communication channel across an IP network at the expense of performance

UDP : the User Datagram Protocol (aka the Unreliable Datagram Protocol) – provides a **lightweight** communication channel across an IP network at the expense of program complexity (i.e., extra effort is needed if reliable channel is needed).

ICMP: the Internet Control Message Protocol – provides a method for handling error conditions, and network control.

Server Identification

- The Internet protocols divide server/service identification into two pieces:
 - an identifier for the *computer* on which a service runs (**IP**)
 - an identifier for a *service* on the computer (**Port Number**)
- Identifier for the computer
 - Unique 32 (or 128) bit identifier known as an Internet Protocol address (IP address)
 - Client must specify the server's IP address
 - Each computer is also assigned a name, and the **Domain Name System** (DNS) is used to translate names into addresses
 - Thus, a user can specify a name such as www.cse.nd.edu rather than an integer address
- Identifier for a service
 - Assigned a unique 16-bit identifier known as a protocol port number (or port number)
 - Email = port number 25, and the Web = port number 80, SSH = port number 22...
 - When a service begins execution...
 - it registers with its local OS by specifying the port number it will use
 - When a client contacts a remote server to request service...
 - the request contains a port number
 - When a request arrives at a server
 - software on the server uses the port number in the request to determine which should handle the request (demultiplexing)
- Summary (with order)



Communication Channels

- The TCP/IP protocol suite provides two communication protocols: **TCP and UDP**
 - TCP is a *high-overhead* protocol. The network code takes care of many error conditions, and provides a reliable communication channel. But due to this overhead, the throughput suffers.
 - UDP is a *low-overhead* protocol. The programmer has to provide the error handling code. The channel is fast, but reliability suffers
- **Message Paradigm (UDP)**
 - network accepts and delivers messages
 - messages are self-contained
 - protocol simply forwards based on the self-contained address
 - N bytes in a msg, N bytes received on the other side
 - spray and pray - no guarantees
 - may be lost, deliver out of order, or duplicated. creates more burden on programmer

- **Stream** Paradigm (TCP)
 - stream: sequence of bytes flow from one application program to another
 - bi-directional: client to server, server to client
 - no boundaries or size limits
 - reliable data delivery
- Comparison

Stream Paradigm (TCP)	Message Paradigm (UDP)
Connection oriented	Connectionless
1-to-1 communication	Many-to-many communication
Sequence of individual bytes	Sequence of individual messages
Arbitrary length transfer	Fixed message size
Used by most applications	Used for multimedia applications
Built on TCP protocol	Built on UDP protocol

Ok, since we have the architecture, now what?


Gotta start developing code to implement the service functions!

We need to understand what TCP/IP provides in respect to an application programming interface (API).

- The underlying feature of TCP/IP programming is the idea of a communication socket.
- The OS provides a bunch of functions and libraries you can call to develop your service code.

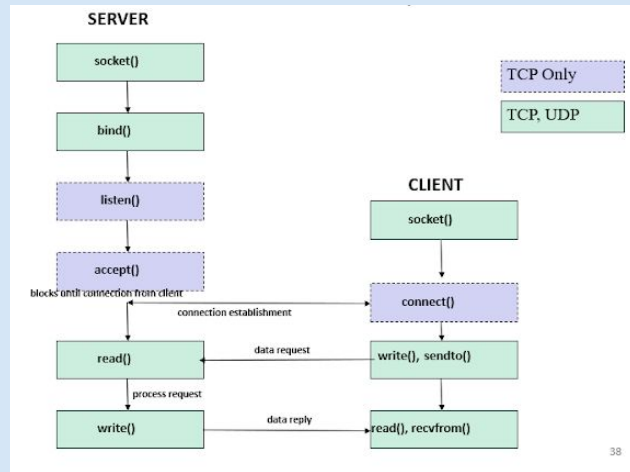
Sockets

- Socket programming
 - Adapter that gets an app onto the network
 - API
 - Depend on the type of socket



- Socket operations
 - Specify local/remote communication endpoints
 - Client: Initiate a connection; Server: Wait for a connection
 - Send/receive data; handle incoming data; terminate connection gracefully; handle error conditions; allocate and release local resources; etc.
- Socket communication
 - kinda similar to **file** I/O... you have open, close, read, write, seek, etc.

- think of socket as the **file descriptor** for network communications, and we can create, read, write and close the socket
- the order of these operations:



(wow tbt old systems days...)

- a nice guide to network programming: Beej's guide to network programming
<http://beej.us/guide/bgnet/>